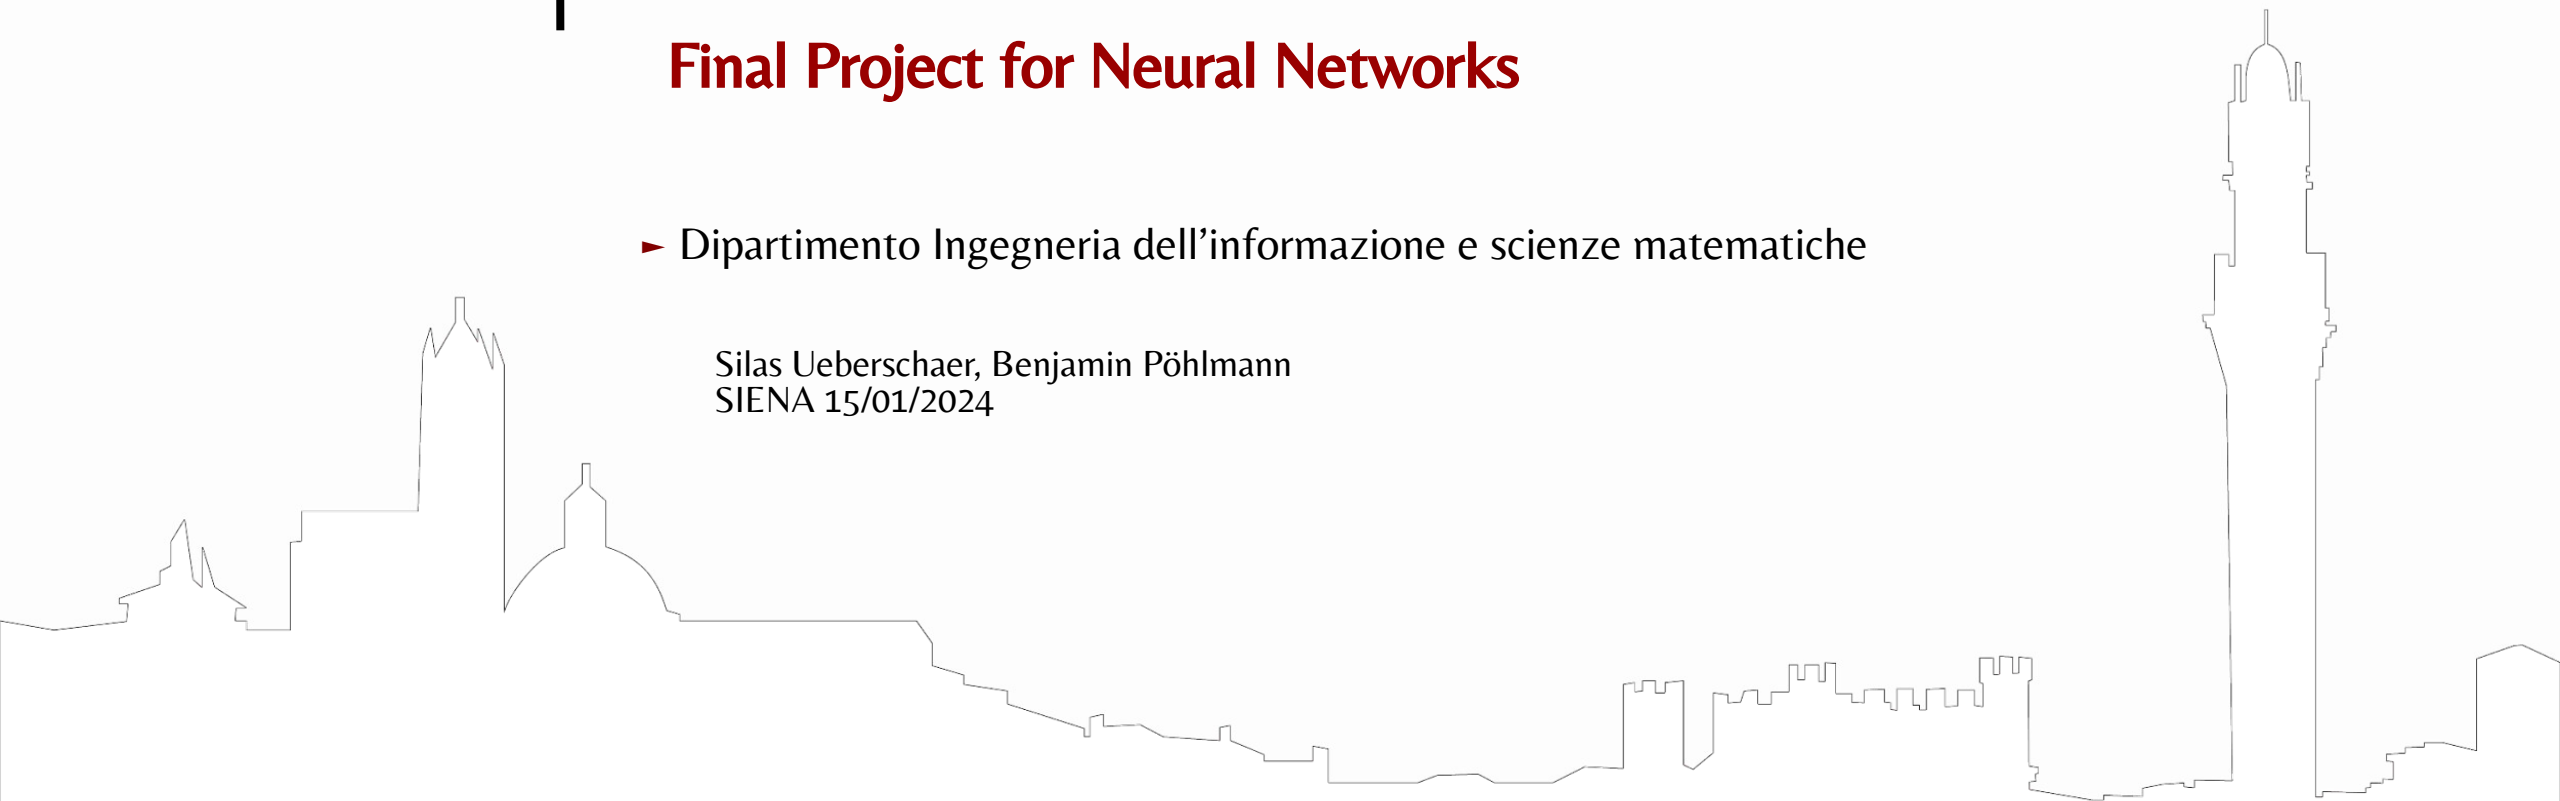


Neural Network for Chess Evaluation

Final Project for Neural Networks

► Dipartimento Ingegneria dell'informazione e scienze matematiche


Silas Ueberschaer, Benjamin Pöhlmann
SIENA 15/01/2024



Goal for the Project

- ▶ **Creating a Neural Network that can predict who is currently winning**
- ▶ **Create our own dataset**
- ▶ **Play against the Neural Network, or let it play against another chess engine (Stockfish)**


Creating our own data



Custom position

- ▶ Using custom start position with reduced amount of pieces to limit possible moves
 - Faster learning
 - Less data
- ▶ Simulating games and labeling them with Stockfish evaluation
 - >0.5 = white losing
 - <0.5 = white winning
 - $=0.5$ = draw
- ▶ Playing a mixture between good moves and random moves to create diverse dataset




Dataset



Dataset Distribution


- ▶ The dataset used for training consists of 70.000 positions
- ▶ Index 0 shows whose turn it is
 - 1 = White
 - 0 = Black
- ▶ Index 1 to 64 is a numerical representation of the chess board
 - 0 = empty, 1 = Pawn, 5 = Rook, 100 = King
- ▶ Index 65 shows the Stockfish evaluation of the position

Architecture



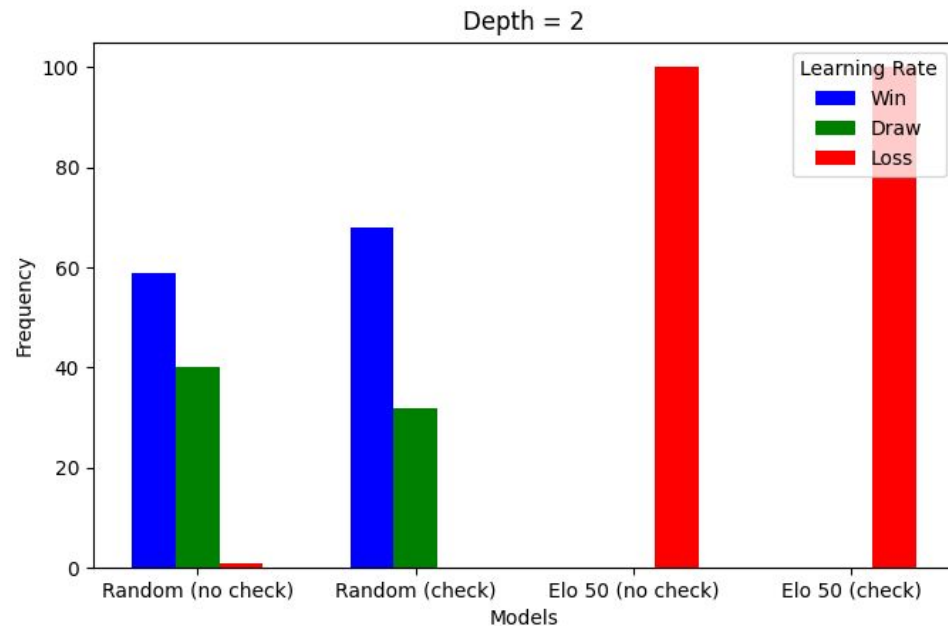
- Simple feedforward neural network with multiple layers and RELU activation function
- To choose the right architecture we trained multiple models with K-Fold Cross-Validation (only on 10.000 positions)
 - $k = 3$
 - learning rate = 0.001/0.01/0.1
 - batch size = 1000/100/10
- Evaluation with MSE, MAE and R2
- Overall best is 3-Layer Neural Network (256,128,64) with learning rate = 0.01 and batch size = 1000

Training



- ▶ Split the Data into Train and Eval
 - 80% Train, 20% Eval
- ▶ MSE Loss function and Adam Optimizer
- ▶ Trained for 400 Epoch
 - learning rate = 0.01
 - batch size = 1000
- ▶ Reached little to no improvement after around 200 epochs on the eval dataset

Playing



- ▶ Plays all moves and evals the board after that move and plays the best move
- ▶ Uses minimax to look multiple moves ahead
- ▶ Possibility to check if position is checkmate and return custom value instead of eval value
 - model not trained to detect checkmate
- ▶ Results (depth=2):
 - wins or draws against random moves
 - loses against stockfish at elo 200

Reinforcement Learning



Reduced number of pieces

- ▶ **Goal: Use Reinforcement Learning to teach a neural network to play chess**
- ▶ **Deep Q Network used**
 - Map a state and action to a Q-Value which is the immediate reward and discounted future rewards
- ▶ **Simplifications used:**
 - Reduce number of pieces (see figure)
 - Assume every piece can theoretically move everywhere
 - Only distinction for reward if move was good or bad but not to what degree

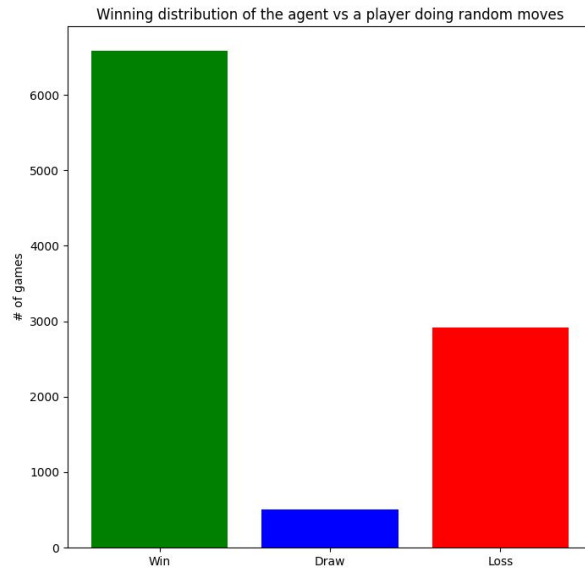
Reinforcement Learning - Setup/Architecture

```
one_hot_mapping = {  
    0: [1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0], # Empty  
    1: [0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0], # White Pawn  
    3: [0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0], # White Bishop  
    4: [0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0], # White Knight  
    5: [0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0], # White Rook  
    1000: [0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0], # White King  
    -1: [0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0], # Black Pawn  
    -3: [0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0], # Black Bishop  
    -4: [0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0], # Black Knight  
    -5: [0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0], # Black Rook  
    -1000: [0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1] # Black King  
}
```

One-Hot-Mapping of chess board

- ▶ **Environment:**
 - Chess board and its rules for legal moves
 - Enemy player (1000 elo)
- ▶ **Agent is white player**
- ▶ **Reward: Evaluation by Stockfish if move increased probability to win, in general:**
 - +1 if increased
 - -1 if decreased
- ▶ **Input is a one-hot-representation of the board for each square on the board**
- ▶ **Output is Q-Value for each action that can theoretically be taken (even if illegal in current board)**

Data creation, training and result



Win distribution of playing 10.000 games of max. 20 turns versus an enemy doing random moves

- ▶ **Play games with limited amount of turns (max. 20 each) to focus on early game**
- ▶ **Using epsilon-greedy algorithm to make mostly random moves**
- ▶ **For each move we save as training example**
 - **Current state (current board)**
 - **Action taken (move we did)**
 - **Reward for our action (change of Stockfish evaluation)**
 - **Next state (board after we and the enemy player moved)**
- ▶ **First 64 elements in output are actions of first pawn, then second pawn,**

Problems and Future Work

- ▶ **Dataset is too little and focuses more on early game**
 - Create more data
- ▶ **Try out more ways to convert chess board**
 - extra neurons for check, checkmate, castle, etc
 - extra neuron that counts all the piece values together
 - ...
- ▶ **The models are still “small” for such a complex topic**
 - use bigger models -> longer training
- ▶ **Instead of using custom position train it on the standard chess position**
 - Needs even more data and bigger models