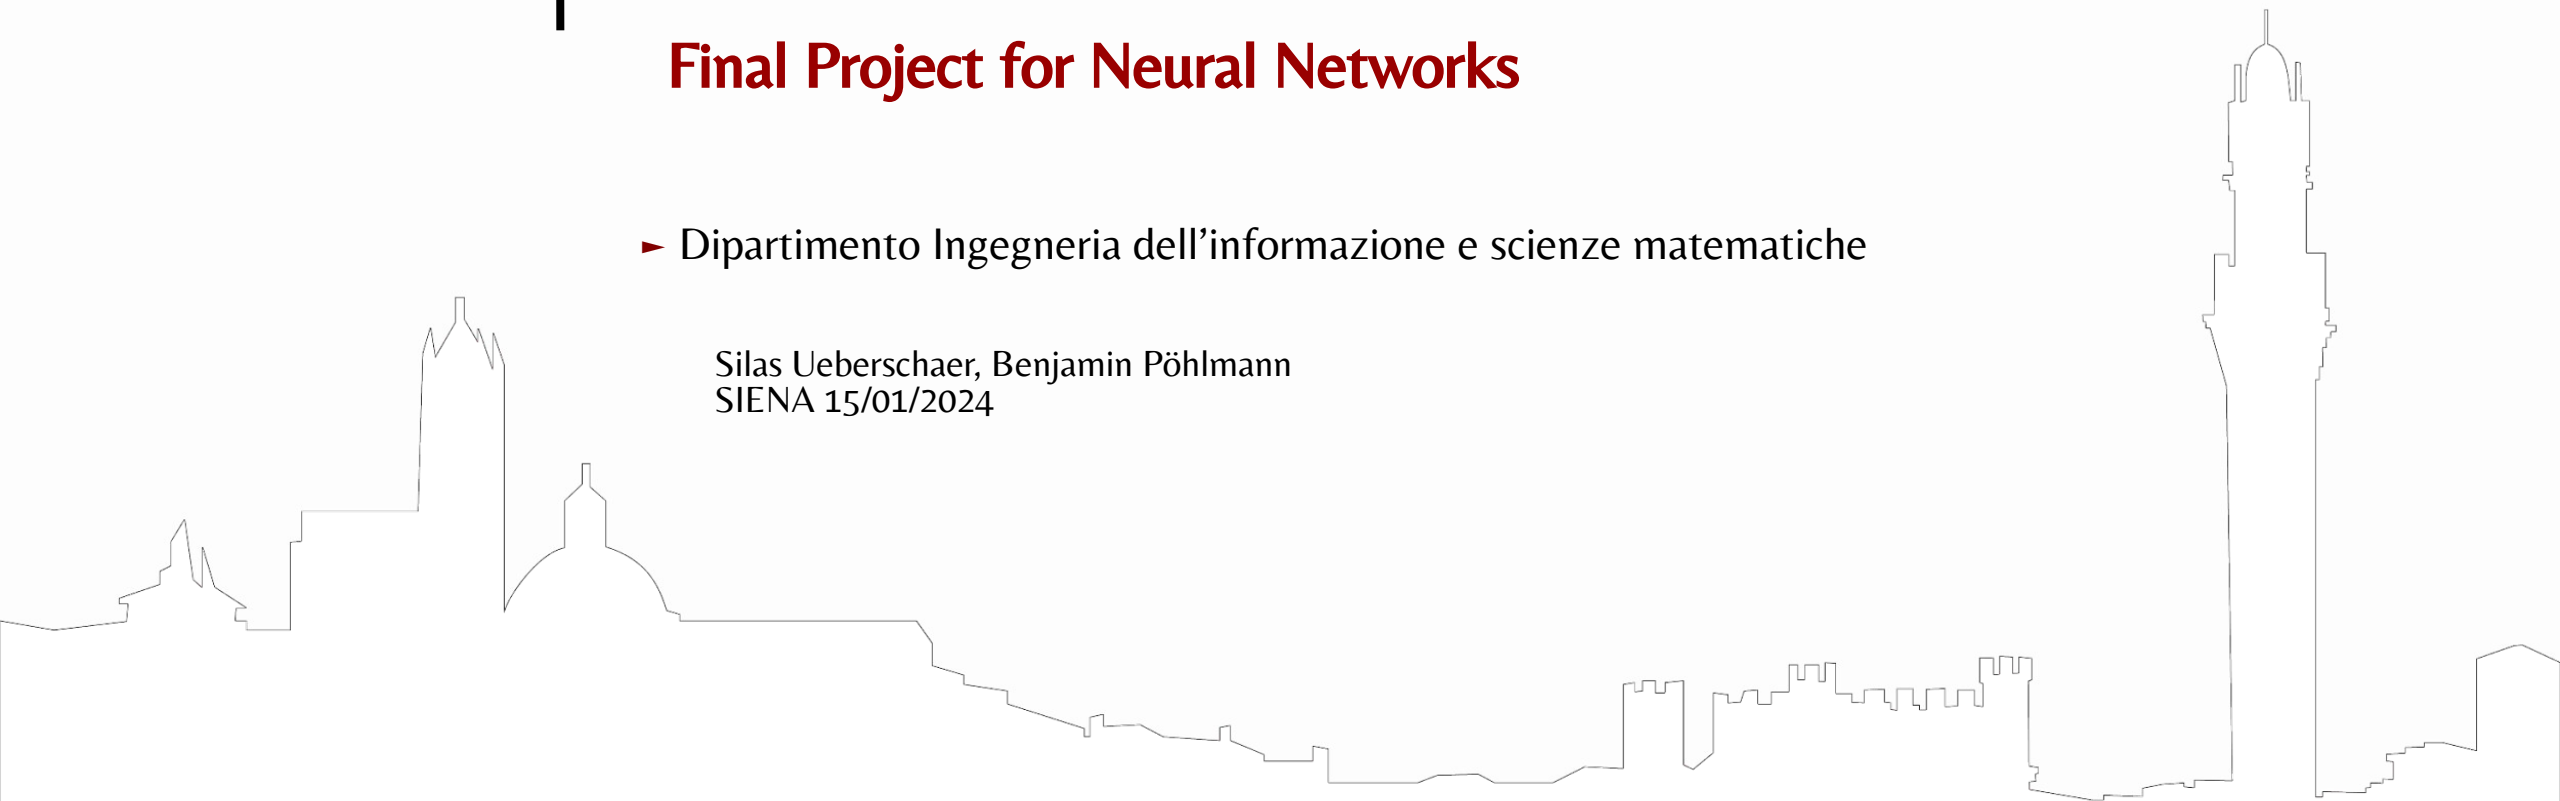


# Neural Network for Chess Evaluation

## Final Project for Neural Networks

► Dipartimento Ingegneria dell'informazione e scienze matematiche


Silas Ueberschaer, Benjamin Pöhlmann  
SIENA 15/01/2024



## Goal for the Project

- ▶ **Creating a Neural Network that can predict who is currently winning**
- ▶ **Create our own dataset**
- ▶ **Play against the Neural Network, or let it play against another chess engine (Stockfish)**

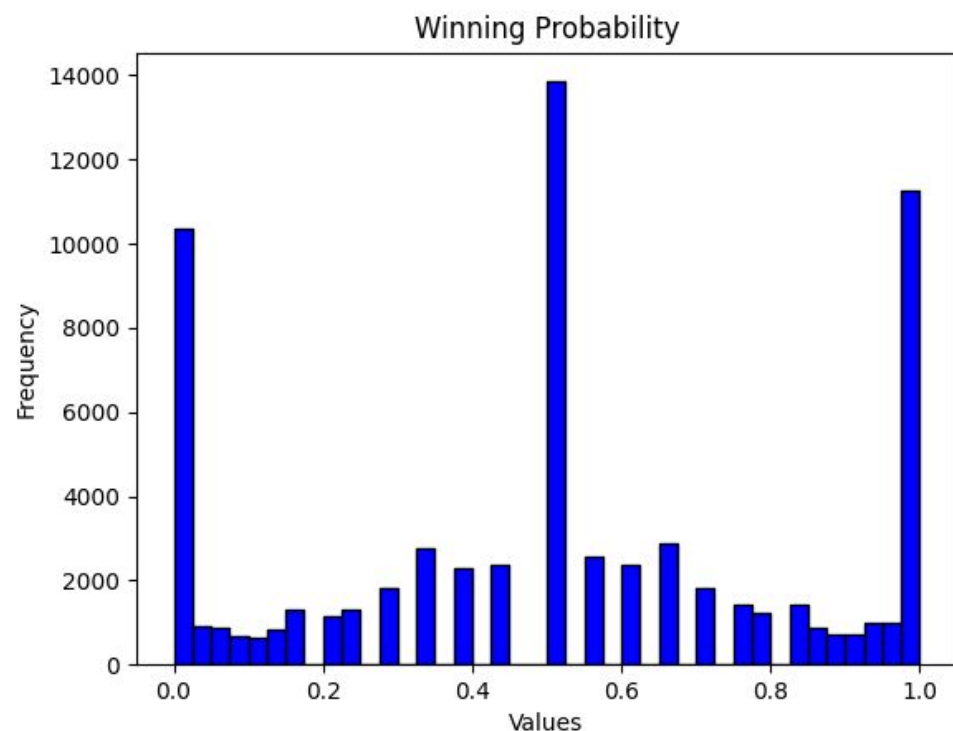
## Creating our own data



Custom position

- ▶ Using custom start position with reduced amount of pieces to limit possible moves
  - Faster learning
  - Less data
- ▶ Simulating games and labeling them with Stockfish evaluation
  - $>0.5$  = white losing
  - $<0.5$  = white winning
  - $=0.5$  = draw
- ▶ Playing a mixture between good moves and random moves to create diverse dataset




# Dataset



Dataset Distribution


- ▶ The dataset used for training consists of 70.000 positions
- ▶ Index 0 shows whose turn it is
  - 1 = White
  - 0 = Black
- ▶ Index 1 to 64 is a numerical representation of the chess board
  - 0 = empty, 1 = Pawn, 5 = Rook, 100 = King
- ▶ Index 65 shows the Stockfish evaluation of the position

# Architecture



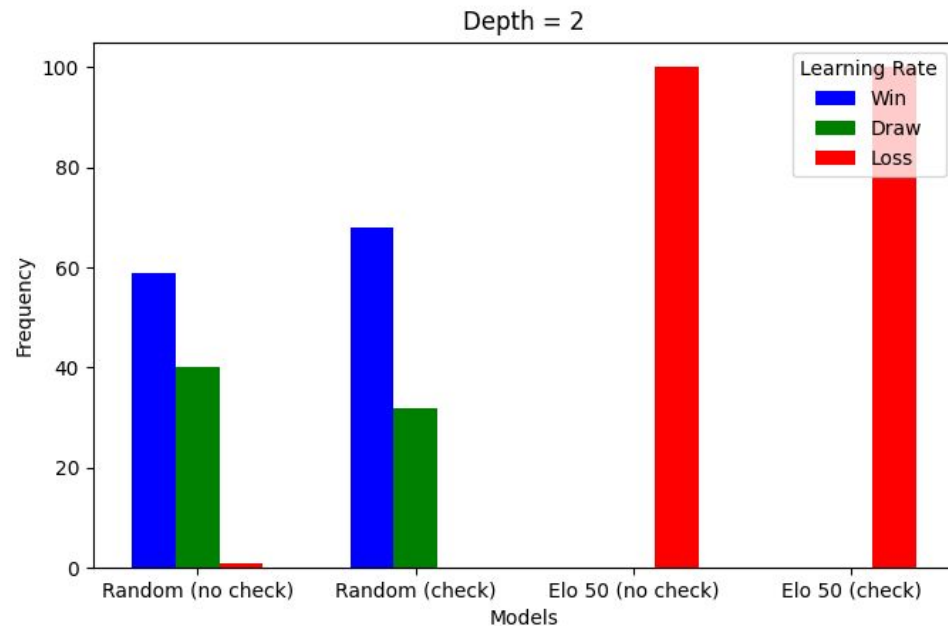
- Simple feedforward neural network with multiple layers and RELU activation function
- To choose the right architecture we trained multiple models with K-Fold Cross-Validation (only on 10.000 positions)
  - $k = 3$
  - learning rate = 0.001/0.01/0.1
  - batch size = 1000/100/10
- Evaluation with MSE, MAE and R2
- Overall best is 3-Layer Neural Network (256,128,64) with learning rate = 0.01 and batch size = 1000

# Training



- ▶ Split the Data into Train and Eval
  - 80% Train, 20% Eval
- ▶ MSE Loss function and Adam Optimizer
- ▶ Trained for 400 Epoch
  - learning rate = 0.01
  - batch size = 1000
- ▶ Reached little to no improvement after around 200 epochs on the eval dataset

# Playing



- ▶ Plays all moves and evals the board after that move and plays the best move
- ▶ Uses minimax to look multiple moves ahead
- ▶ Possibility to check if position is checkmate and return custom value instead of eval value
  - model not trained to detect checkmate
- ▶ Results (depth=2):
  - wins or draws against random moves
  - loses against stockfish at elo 200

# Reinforcement Learning



Reduced number of pieces

- ▶ **Goal: Use Reinforcement Learning to teach a neural network to play chess**
- ▶ **Deep Q Network used**
  - Map a state and action to a Q-Value which is the immediate reward and discounted future rewards
- ▶ **Simplifications used:**
  - Reduce number of pieces (see figure)
  - Assume every piece can theoretically move everywhere
  - Only distinction for reward if move was good or bad but not to what degree



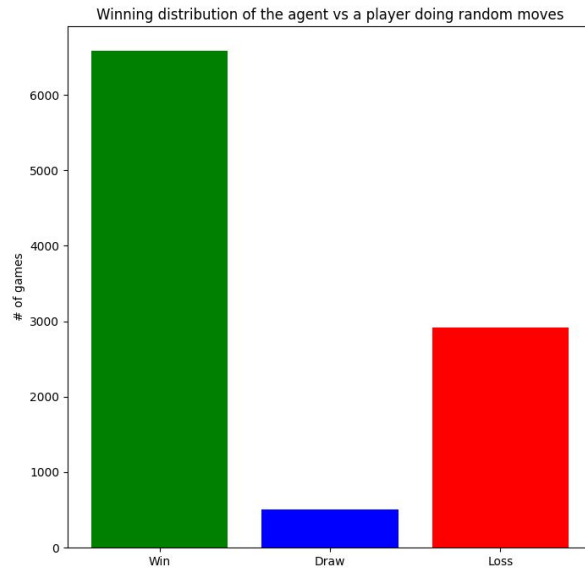
# Reinforcement Learning - Setup/Architecture

```
one_hot_mapping = {  
  0: [1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0], # Empty  
  1: [0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0], # White Pawn  
  3: [0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0], # White Bishop  
  4: [0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0], # White Knight  
  5: [0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0], # White Rook  
  1000: [0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0], # White King  
  -1: [0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0], # Black Pawn  
  -3: [0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0], # Black Bishop  
  -4: [0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0], # Black Knight  
  -5: [0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0], # Black Rook  
  -1000: [0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1] # Black King  
}
```

## One-Hot-Mapping of chess board

- ▶ **Environment:**
  - Chess board and its rules for legal moves
  - Enemy player (1000 elo)
- ▶ **Agent is white player**
- ▶ **Reward: Evaluation by Stockfish if move increased probability to win, in general:**
  - +1 if increased
  - -1 if decreased
- ▶ **Input is a one-hot-representation of the board for each square on the board**
- ▶ **Output is Q-Value for each action that can theoretically be taken (even if illegal in current board)**

## Data creation, training and result



**Win distribution of playing 10.000 games of max. 20 turns versus an enemy doing random moves**

- ▶ **Play games with limited amount of turns (max. 20 each) to focus on early game**
- ▶ **Using epsilon-greedy algorithm to make mostly random moves**
- ▶ **For each move we save as training example**
  - **Current state (current board)**
  - **Action taken (move we did)**
  - **Reward for our action (change of Stockfish evaluation)**
  - **Next state (board after we and the enemy player moved)**
- ▶ **First 64 elements in output are actions of first pawn, then second pawn, ....**

## Problems and Future Work

- ▶ **Dataset is too little and focuses more on early game**
  - **Create more data**
- ▶ **The models are still “small” for such a complex topic**
  - **use bigger models -> longer training**
- ▶ **Instead of using custom position train it on the standard chess position**
  - **Needs even more data and bigger models**